



Amdahl's law: theoretical speedup in latency of the execution of a task at fixed workload.
Amdahl's Law
u f: Parallelizable fraction of a program
u N: Number of processors
Speedup = $\frac{1}{1-f + \frac{f}{N}}$
Serial bottleneck of Amdahl's Law: Maximum speedup (1/(1-f)) limited by serial portion (1 - f)
Parallel portion (f) is usually not perfectly parallel: Synchronization overhead (e.g., updates to shared data), Load imbalance overhead (imperfect parallelization), Resource sharing (contention among processors)
Roofline model:Theoretical performance bound of your application running on your machine.
Computing regime: Latency-limited -> throughput-limited
Roofline Model's Two Perspectives:
1, Target processor's perspective: Showing inherent hardware limitations (or bound), in term of compute and memory
2, Compute kernel's perspective : Showing the priority of optimizations for a given compute kernel running on a given processor
Arithmetic intensity (AI): AI = Total Flops / Total Memory Bytes, describes the characteristics of a compute kernel running on a given processor
Large AI - Compute-bound, Small AI - Memory-bound
Roofline model's 3 Steps:
1, Machine characterization: Memory bandwidth, Peak compute; 2, Application Characterization: Arithmetic intensity; 3, Application execution monitoring: Real Throughput
 $\text{Attainable Flops} = \min(\text{peak Flops/s}, \text{AI} * \text{peak GB/s})$
★ AI = $\frac{\text{peak Flops/s}}{\text{peak GB/s}}$

斜率:cache>HBM>DEAM,平台: peak flops>vec>no vectorization
Little's Law: $L = \lambda * W$ (buffer size = throughput * latency)
The long-term average number L of customers in a stationary system is equal to the long-term average effective arrival rate λ multiplied by the average time W that a customer spends in the system.
Von Neumann model
Components: memory(stores the program and data), processing unit, input, output, control unit(controls the order in which instructions are carried out)

Memory stores programs and data, bit logically grouped into bytes and words(8/16/32 bits)
Address space: total number of uniquely identifiable locations in memory.
Addressability/how many bits are stored in each location/address (8-bit(byte)/word-addressable)
PU(processing unit) performs the actual computations, consists of ALU and temporary storage(register)
ALU execute computation and logic operations, combine a variety of arithmetic and logic operation into a single unit,perform only one function at a time.
Registers:ensure fast access to values to be processed in ALU, one reg contain one word.
Register set/file:set of registers manipulated by instructions.MIPS 32 general purpose registers.
Input and output enables information to get into and out of a computer.
Control unit:conductor of an orchestra.conduct execution,track instruction being process(IR) and to be processed(PC & IP instruction pointer).
Key properties of Von Neumann Model: stored program, sequential instruction processing.
ISA: Interface between software command and hardware carry-out.specify memory organization(address space/ addressability ,word/byte addressable) register set(32 in MIPS), instruction set(opcodes, datatype, addressing mode(immediate or literal, register, memory addressing modes(PC-relative, pseudo-direct, base-offset)),more addressing modes create more work for compiler and microarchitect.
ABI: application binary interface, between two binary program modules.
Semantic gap:语义鸿沟, 存在于高层次语义和低层次操作之间。
Opcode: what the instruction does;
Operands: who the instruction is to do with.source/destination ~
Type of opcodes: operate(R/I/F), data movement, control(conditional branches and unconditional jumps)
Complex instructions:~denser encoding ,+simpler compiler ,+larger chunks of work ,+complex hardware
Large number of registers: better register allocation but larger instruction size and larger register file size.
Instruction cycle: the process of executing an instruction.
IF-ID-EXE-MEM-WB. LDR no EXE, R/I no MEM,SW no WB, branch no MEM,WB, jump no ID,EX,MEM,WB
Single-cycle CPU:AS/architecture state)-process-AS'(state,next-state logic), combinational logic,无中间状态
Microarchitecture implement how AS transformed to AS'.
MIPS state element: PC, instruction memory/register file, data memory.
Clock time determined by slowest instruction.
Multi-cycle CPU
Decrease clock cycle time; each instruction take as many clock cycle as it needs to take.
AS->AS+MS1->AS+MS2->AS+MS3->AS'
Benefit:reduce critical path; bread and butter design (优化常用指令); balanced design (硬件资源复用)
Execution time: CPI(>1) * clock cycle time(short)



Pipeline CPU: more concurrency and higher instruction throughput(1/CPI).
Ideal assumption: repetition of identical/independent operations; uniformly partitionable suboperations.
Lec2: pipeline hazard+reorder buffer
Pipeline hazards: prevent instruction from executing next stage.
Structural hazards:limited hardware resources-replicating hardware resources.
Data hazards: flow dependence(RAW, 顺序依赖), output dependence(WAW), anti dependence(WAR)
RAW: hardware pipeline stall; software pipeline stall(nops); data forward/bypass
Control hazard: data dependence on IP/PC.
Reorder buffer
Exceptions: internal problems,handle when detected. Interrupts: external events,handle when convenient.
ROB: complete instructions out of order,reorder before writing result to AS.




ROBs need to:
1, correctly reorder instructions back into the program order
2, update the architectural state with the instruction's result(s), if instruction can retire without any issues
3, handle exception/interrupt precisely, if needs to be handled before retiring the instruction
4, use valid bits to keep track of readiness of the result, find out if the instruction has completed execution
Lec3: Tomasula
Solution: out-of-order dispatch
Reservation station: move dependent instructions out of the way of independent ones.instructions dispatched in dataflow order. benefit: latency tolerance.
Two humps in modern pipeline: reservation stations(enable in-order issue and out-of-order dispatch), reorder buffer(enable OoO completion, in-order commitment)
1. register renaming: 2. insert into reservation station; 3.broadcast tag when value produced and compare; 4.all values ready, dispatch into FU(functional unit).
tomasula 三大组件: 寄存器重命名表(RAT,register alias table), RS, 公共数据总线(CDB)
Dataflow graph limited to instruction window(all decoded not retired instructions)
Superscalar+SIMD+multi-core
Superscalar execution: N-wide superscalar->N instructions per cycle
Need to add hardware;hardware check dependence between concurrently-fetched instruction. +higher throughput -higher complexity for dependence checking -more hardware needed
Vector Insns
SISD: Single instruction operates on single data element
SIMD: Single instruction operates on multiple data elements (Array processor ,Vector processor)
MISD: Multiple instructions operate on single data element (Closest form: systolic array processor, streaming processor)
MIMD: Multiple instructions operate on multiple data elements (multiple instruction streams) (Multiprocessor, Multithreaded processor)
For SIMD,Memory (bandwidth) can easily become a bottleneck, especially if compute/memory operation balance is not maintained or data is not mapped appropriately to memory banks
增加向量寄存器堆, multiple bank (内存分为多个块)
scalar 和 SIMD 提高 peak flops
Fine-grained multithreading: Hardware has multiple thread contexts (PC=registers). Each cycle, fetch engine fetches from a different thread.
无需检查指令间依赖和分支预测逻辑, 提高吞吐量; 硬件复杂度高, 单线程性能下降; 资源竞争; 仍需依赖检查, 如访存.
可以使实际性能上移, 更接近理论峰值. 增加带宽压力, cache命中率下降
Multi-core processors:Put multiple cores on the same die.
Alternative: bigger, more powerful single core: 对程序员透明, 提高单线程性能; 功耗高, 收益递减, 设计复杂.
多核: 简单, 能效高, 频率容易提升, 吞吐量高; 需要并行编程; 单线程性能低; 引脚带宽限制.
Technology push: instruction issue queue; large, multi-ported register files
成倍提高 peak flops, 加剧方寸受限

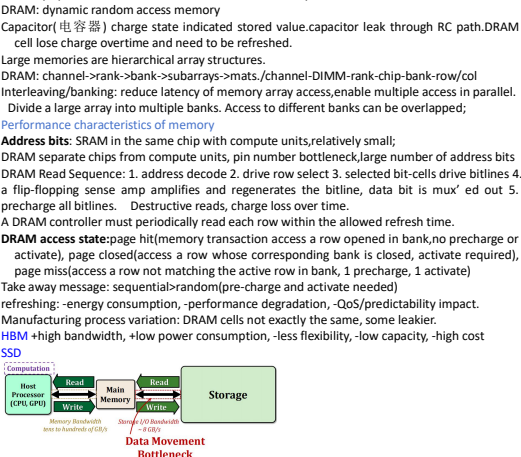
Memory of physical memory: zero latency,infinite capacity,infinite bandwidth, zero cost
Bigger is slower; faster is more expensive;higher bandwidth is more expensive.

Flip-Flops (~'K'):Very fast, parallel access, Very expensive (one bit costs tens of transistors)
Static RAM (~'M'):Relatively fast,one data word at a time ,Expensive
Dynamic RAM (~'G'):Slow, one data word at a time, reading destroys content (refresh), needs special process for manufacturing ,Cheap (one bit costs only one transistor plus one capacitor)

Flash Memory (~'T'):Much slower, access takes a long time, non-volatile ;Very cheap
Array organization of memory: to efficiently store large amounts of data: memory array(store data),address selection logic(one row of array), readout circuitry(read data out)
Two-dimensional array of bits: each bit cell store one bit.N address bits and M data bits.2^n rows, M columns.depth:number of rows(number of words), width:number of cols(size of word), array size: depth * width=2^n * M
Bitline: Storage nodes in one column connected to one bitline
Wordline: Address decoder activates only one wordline, content of one line of storage available at output



SRAM : a bit of static random access memory
Access transistors connect the bit storage to the bitline; access controlled by the wordline
Goal: buffering data on chip to reduce external memory traffic, high performance ,low capacity.
Used in cache in CPU, shared memory in GPU, in-chip buffer in AI accelerator.
Memory banking: memory divided into banks that can be accessed independently; share address and data buses; start and complete one bank access per cycle.
Read Sequence: 1. address decode 2. drive row select 3. selected bit-cells drive bitlines (entire row is read together) 4. differential sensing and column select (data is ready) 5. precharge all bitlines (for next read or write)
Access latency dominated by steps 2 and 3. Cycling time dominated by steps 2, 3 and 5
step 2 proportional to 2^n,m, step 3 and 5 proportional to 2^n
DRAM
Key components of computing system: computation, communication, storage/memory
Computation is bottlenecked by memory(important workloads data intensive, data increasing)
Performance/energy/reliability perspective: memory as bottleneck.
High level abstraction: map virtual memory addresses to physical memory
Memory optimization aim at size, bandwidth, not latency.
DRAM: dynamic random access memory
Capacitor (电容器) charge state indicated stored value.capacitor leak through RC path.DRAM cell lose charge overtime and need to be refreshed.
Large memories are hierarchical array structures.
DRAM: channel->rank->bank->subarrays->mats ,channel-DIMM-rank-chip-bank-row/col
Interleaving/banking: reduce latency of memory array access,enable multiple access in parallel.
Divide a large array into multiple banks. Access to different banks can be overlapped;
Performance characteristics of memory
Address bits: SRAM in the same chip with compute units,relatively small;
DRAM separate chips from compute units, pin number bottleneck,large number of address bits
DRAM Read Sequence: 1. address decode 2. drive row select 3. selected bit-cells drive bitlines 4. a flip-flopping sense amp amplifies and regenerates the bitline, data bit is mux' ed out 5. precharge all bitlines. Destructive reads, charge loss over time.
A DRAM controller must periodically read each row within the allowed refresh time.
DRAM access state:page hit(memory transaction access a row opened in bank,no precharge or activate), page closed(access a row whose corresponding bank is closed, activate required), page miss(access a row not matching the active row in bank, 1 precharge, 1 activate)
Take away message: sequential+random(pre-charge and activate needed)
refreshing: -energy consumption, -performance degradation, -QoS/predictability impact.
Manufacturing process variation: DRAM cells not exactly the same, some leakier.
HBM -high bandwidth, +low power consumption, -less flexibility, -low capacity, -high cost
SSD



Phase change material exists in amorphous(low optical reflexivity, high electrical resistivity) and crystalline(high~, low~). PCM is resistive memory: high resistance(0), low~(1)
Graphics processing units
CPU: few complex cores, large cache, large and slow memory
GPU:lots of simple cores, small cache, small and fast memory
GPU computing: computation offloaded to GPU. CPU-GPU data transfer(1)->GPU kernel execution(2)->GPU-CPU data transfer(3)
GPU are SIMD engines underneath, yet programming done using threads
Programming model(software):how programmer express code,like data parallel, dataflow,...
Hardware execution model: how hardware execute the code, OoO, vector processor,...
Programming model
Sequential(SISD):executed on:pipelined, OoO execution, superscalar/VLIW processor
Data Parallel(SIMD): generate SIMD instruction
Multithread(SPMD, single program multiple data): generate thread to execute each iteration.
SPMD: multiple instruction stream execute same program, different data and control-flow path
CUDA(SPMD example):grid, thread block, thread
CUDA program structure:
1.Function prototypes:float serialFunction(~~); __global__ void kernel(~~);
2.main()
Memory allocation :cudaMalloc((void**) &in, n*bytes);
Memory copy :cudaMemcpy(d_in, h_in, n*bytes, cudaMemcpyHostToDevice);
Kernel launch :kernel<<< #blocks, #threads >>>(args);
Memory deallocation :cudaFree(d_in);
Explicit synchronization :cudaDeviceSynchronize();
3.Kernel_ __global__ void kernel(type args, ~~)
blockDim: block dimension, blockIdx: block id within grid , threadIdx: thread index within block
SIMT(hardware) & warp(software)
SIMTsingle instruction multiple thread; warp: basic execution unit with 32 consecutive threads, thread block divided into warps for SIMT execution.can reduce GPU scheduling overhead
SIMT: + treat each thread separately, +group threads into warps flexibly
GPU optimization
GPU memory: optimize with multithreading, memory coalescing, shared memory
FGMT hide long latency operations, occupancy=active warps/maximum number of warps
Memory coalescing: thread in same warp access consecutive memory location in same burst,

combine and serve by one burst.make sure that concurrent threads access nearby memory location. Peak bandwidth utilization when all threads in warp access one cache line.
Memory divergence: location not in same burst, multiple transactions needed, longer time
Shared memory: interleaved/banked memory, bank-address%12(NVIDIA)
Bank conflicts only possible within a warp. Due to strided access.
Optimization: padding, randomized mapping, hash functions
Data reuse: tiling: divide input into tiles, load L_size chunks into shared memory, compute.
Loading/compute amount: (L_size+2)^2/L_size^2 per thread.
SIMT efficiency
Branch divergence occur when threads inside warps branch to different execution paths.
Atomic operation: on shared/global memory,perform read-modify-write operation atomically.
Atomic conflict=serialized updates, prevent data races when two thread update same location.
Estimates $t_e + \frac{t_r}{\#Streams}$ $t_e + \frac{t_r}{\#Streams}$ $t_e + \frac{t_r}{\#Streams}$
 $t_e > t_r$ (dominant kernel) $t_e > t_r$ (dominant kernel) $t_e > t_r$ (dominant kernel)
Dynamic warp formation/merging: merge threads executing same instruction.

Generic Term	NVIDIA Term	AMD Term	Comments
Vector length	Warp size	Wavefront size	Number of threads that run in parallel (lock-step) on a SIMD functional unit
Pipelined functional unit / scalar pipeline	Streaming processor / CUDA core	-	Functional unit that executes instructions for one CPU thread
SIMD functional unit / SIMD pipeline	Group of N streaming processors (e.g., N=8 in GTX 285, N=16 in Fermi)	Vector ALU	SIMD functional unit that executes instructions for an entire warp
GPU core	Streaming multiprocessor	Compute unit	It contains one or more warp schedulers and one or several SIMD pipelines

Cache:Have multiple levels of storage, most data processor needs kept in faster levels.
Memory hierarchy:regfile-> L1 cache-> L2 cache->L3 cache->main memory->swap disk
Cache feasibility: temporal locality, spatial locality
Recently accessed data will be again accessed in the near future.
Nearby memory locations will be accessed soon.
 $T_r = h_r * t_r + m_r * (t_r + T_{w,r})$ $T_r = t_r + m_r * T_{w,r}$ $T_r = t_r + m_r * T_{w,r}$ $T_r = t_r + m_r * T_{w,r}$
r_l: hit rate, mr_l: miss rate
Cache: an automatically-managed memory structure
Blocks: main memory logically divided into fixed-size chunks.cache house limited number of it.
Each block address maps to a potential location in cache.
Cache hit rate = hits/(hits+misses)=hits/accesses
AMAT(average memory access time)=(hit-rate*hit-latency)+(miss-rate*miss-latency)
Cache access: 1) index into the tag and data stores with index bits in address; 2) checks valid bit in tag store; 3) compares tag bits in address with the stored tag in tag store; 4) if a block is in the cache (cache hit), the stored tag should be valid and match the tag of the block, and read out data.
Cache block(line):unit of storage in cache.
Placement: where and how to place/find a block in cache?
Directed-mapped: a chunk go to only one cache block.+easy to implement.-low hit rate
Fully associative: a chunk go to any cache block.+highly utilization.-difficult to implement.
Set-associative: a chunk go to N cache blocks in N-way set-associative cache.+accommodate conflict better.more complex, slower access, larger tag store.
Degree of associativity: how many blocks can map to the same index.high associativity->high hit rate, slow cache access time, expensive hardware.
Replacement: what data to remove to make room in cache?
Any invalid block first. If all valid:random, FIFO, least recently used, hybrid ~, optimal~
LRU: need to keep track of access ordering of blocks,use approximation.
Pseudo LRU:PLRU replacement way selection, PLRU bits updating rule.
Set thrashing: program working set larger than set associativity, random better.
Belady's OPT:replace the block going to be referenced furthest.
Granularity of management: large or small blocks? Subblocks?
Write policy: what do we do about writes?
1. store issn->cache:write-allocate policy(memory, default, allocate the cache line), write-no-allocate policy(Pcie/I/O,write directly to memory without allocation in cache)
2. Cache->memory: write-back policy(default, write to cache until cache kick cache block out), write-through policy(streaming write instruction, write to cache and memory right away)
Instructions/data: do we treat them separately?
Unified cache:+dynamic sharing of cache space.-instructions and data thrash each other.-I and D crossed in different phases in pipeline.
Cache miss:compulsory miss(first reference, prefetching) capacity miss(cache too small,utilize space, software management) conflict miss(other more associativity(victim cache, better, randomized indexing))
Improve cache performance: reduce miss rate, reduce latency/cost, reduce hit latency/cost
Capacity (C): the number of data bytes a cache stores
Block size (b): bytes of data brought into cache at once
Number of blocks (B = C/b): number of blocks in cache: B = C/b
Degree of associativity (N): number of blocks in a set
Number of sets (S = B/N): each memory address maps to exactly one cache set

Cache in multi-core CPU
Private cache:cache belong to one core.
Shared cache: cache shared by multiple cores.
Resource sharing: allow multiple contexts to use hardware resource.+improve utilization/efficiency(throughput);+reduce communication latency;+compatible with shared memory programming model;-reduce some thread's performance;-eliminate performance isolation
Shared cache between cores: +high effective capacity; +dynamic partitioning of available cache space;+easier to maintain coherence; -slower access; -conflict miss; -hard to guarantee a minimum level of service.
Coherence: ordering of operations from different processors to the same memory location.
Consistency: ordering of all memory operations from different processors.

